

[About Keras](#)[Getting started](#)[Developer guides](#)[Keras API reference](#)

Code examples

Computer Vision

Natural Language Processing

Structured Data

Timeseries

Audio Data

Generative Deep Learning

Reinforcement Learning

Graph Data

Quick Keras Recipes

[Why choose Keras?](#)[Community & governance](#)[Contributing to Keras](#)[KerasTuner](#)» [Code examples](#) / [Computer Vision](#) / Simple MNIST convnet

Simple MNIST convnet

Author: [fchollet](#)**Date created:** 2015/06/19**Last modified:** 2020/04/21**Description:** A simple convnet that achieves ~99% test accuracy on MNIST.[View in Colab](#) • [GitHub source](#)

Setup

```
import numpy as np
from tensorflow import keras
from tensorflow.keras import layers
```

Prepare the data

```
# Model / data parameters
num_classes = 10
input_shape = (28, 28, 1)

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# Scale images to the [0, 1] range
x_train = x_train.astype("float32") / 255
x_test = x_test.astype("float32") / 255
# Make sure images have shape (28, 28, 1)
x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
print("x_train shape:", x_train.shape)
print(x_train.shape[0], "train samples")
print(x_test.shape[0], "test samples")

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

Build the model

```
model = keras.Sequential([
    keras.Input(shape=input_shape),
    layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Flatten(),
    layers.Dropout(0.5),
    layers.Dense(num_classes, activation="softmax"),
])

model.summary()
```

Model: "sequential"		
Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dropout (Dropout)	(None, 1600)	0
dense (Dense)	(None, 10)	16010
=====		
Total params: 34,826		
Trainable params: 34,826		
Non-trainable params: 0		

Train the model

```
batch_size = 128
epochs = 15

model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])

model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, validation_split=0.1)
```

```
Epoch 1/15
422/422 [=====] - 13s 29ms/step - loss: 0.7840 - accuracy: 0.7643 -
val_loss: 0.0780 - val_accuracy: 0.9780
Epoch 2/15
422/422 [=====] - 13s 31ms/step - loss: 0.1199 - accuracy: 0.9639 -
val_loss: 0.0559 - val_accuracy: 0.9843
Epoch 3/15
422/422 [=====] - 14s 33ms/step - loss: 0.0845 - accuracy: 0.9737 -
val_loss: 0.0469 - val_accuracy: 0.9877
Epoch 4/15
422/422 [=====] - 14s 33ms/step - loss: 0.0762 - accuracy: 0.9756 -
val_loss: 0.0398 - val_accuracy: 0.9895
Epoch 5/15
422/422 [=====] - 15s 35ms/step - loss: 0.0621 - accuracy: 0.9812 -
val_loss: 0.0378 - val_accuracy: 0.9890
Epoch 6/15
422/422 [=====] - 17s 40ms/step - loss: 0.0547 - accuracy: 0.9825 -
val_loss: 0.0360 - val_accuracy: 0.9910
Epoch 7/15
422/422 [=====] - 17s 41ms/step - loss: 0.0497 - accuracy: 0.9840 -
val_loss: 0.0311 - val_accuracy: 0.9920
Epoch 8/15
422/422 [=====] - 16s 39ms/step - loss: 0.0443 - accuracy: 0.9862 -
val_loss: 0.0346 - val_accuracy: 0.9910
Epoch 9/15
422/422 [=====] - 17s 39ms/step - loss: 0.0436 - accuracy: 0.9860 -
val_loss: 0.0325 - val_accuracy: 0.9915
Epoch 10/15
422/422 [=====] - 16s 38ms/step - loss: 0.0407 - accuracy: 0.9865 -
val_loss: 0.0301 - val_accuracy: 0.9920
Epoch 11/15
422/422 [=====] - 16s 37ms/step - loss: 0.0406 - accuracy: 0.9874 -
val_loss: 0.0303 - val_accuracy: 0.9920
Epoch 12/15
237/422 [=====>.....] - ETA: 7s - loss: 0.0398 - accuracy: 0.9877
```

Evaluate the trained model

```
score = model.evaluate(x_test, y_test, verbose=0)
print("Test loss:", score[0])
print("Test accuracy:", score[1])
```

```
Test loss: 0.023950600996613503
Test accuracy: 0.9922000169754028
```